

# Ensemble\_Learning\_Project

May 15, 2022

- 0.1 We are provided with the NYC Taxi Trip Dataset. This dataset contains information about the taxi trips that took place in different parts of New York City and how much time did that trip take to complete.
- 0.2 In this project we have the following tasks as follows:-
  - 0.2.1 1. Build a K-Nearest neighbours model for the given dataset and find the best value of K.
  - 0.2.2 2. Build a Linear model for the given dataset with regularisation.
  - 0.2.3 3. Build a Random Forest model for the given dataset.
  - 0.2.4 4. Build a Gradient Boosting model for the given dataset.
  - 0.2.5 5. Combine all the models above using the averaging technique to generate the final predictions.

```
[1]: import pandas as pd
import numpy as np

from sklearn.preprocessing import MinMaxScaler, LabelEncoder
import geopy.distance

from sklearn.model_selection import train_test_split

from sklearn.neighbors import KNeighborsRegressor
from sklearn.linear_model import LinearRegression
from sklearn.linear_model import Ridge
from sklearn.ensemble import RandomForestRegressor
from xgboost import XGBRegressor

from sklearn.metrics import mean_squared_error
from math import sqrt
from sklearn.metrics import r2_score

import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline

import datetime as dt
```

```
import warnings
warnings.filterwarnings('ignore')
```

```
[2]: data = pd.read_csv('nyc_taxi_trip_duration_dataset.csv')
```

### 0.3 Performing Exploratory Data Analysis and Feature Engineering

#### 0.4 Dataset Variables

##### 0.4.1 Dataset Variables are as follows :

vendor\_id - code indicating provider associated with the trip record

id - unique identifier for each trip

pickup\_datetime - date and time when the trip started

dropoff\_datetime - date and time when the trip ended

passenger\_count - number of passengers present during the trip

pickup\_longitude - the longitude where the trip was started

pickup\_latitude - the latitude where the trip was started

dropoff\_longitude - the longitude where the trip was ended

dropoff\_latitude - the latitude where the trip was ended

store\_and\_fwd\_flag - The flag indicates that whether the connection during the trip with vehicle and vendor was lost due to any technical issue or not

trip\_duration - duration of trip in seconds

```
[3]: data.head()
```

```
[3]:
```

	id	vendor_id	pickup_datetime	dropoff_datetime	\
0	id1080784	2	2016-02-29 16:40:21	2016-02-29 16:47:01	
1	id0889885	1	2016-03-11 23:35:37	2016-03-11 23:53:57	
2	id0857912	2	2016-02-21 17:59:33	2016-02-21 18:26:48	
3	id3744273	2	2016-01-05 09:44:31	2016-01-05 10:03:32	
4	id0232939	1	2016-02-17 06:42:23	2016-02-17 06:56:31	

	passenger_count	pickup_longitude	pickup_latitude	dropoff_longitude	\
0	1	-73.953918	40.778873	-73.963875	
1	2	-73.988312	40.731743	-73.994751	
2	2	-73.997314	40.721458	-73.948029	
3	6	-73.961670	40.759720	-73.956779	
4	1	-74.017120	40.708469	-73.988182	

	dropoff_latitude	store_and_fwd_flag	trip_duration
0	40.771164	N	400
1	40.694931	N	1100
2	40.774918	N	1635
3	40.780628	N	1141
4	40.740631	N	848

```
[4]: data.tail()
```

```
[4]:
```

	id	vendor_id	pickup_datetime	dropoff_datetime	\
729317	id3905982	2	2016-05-21 13:29:38	2016-05-21 13:34:34	
729318	id0102861	1	2016-02-22 00:43:11	2016-02-22 00:48:26	
729319	id0439699	1	2016-04-15 18:56:48	2016-04-15 19:08:01	
729320	id2078912	1	2016-06-19 09:50:47	2016-06-19 09:58:14	
729321	id1053441	2	2016-01-01 17:24:16	2016-01-01 17:44:40	

	passenger_count	pickup_longitude	pickup_latitude	dropoff_longitude	\
729317	2	-73.965919	40.789780	-73.952637	
729318	1	-73.996666	40.737434	-74.001320	
729319	1	-73.997849	40.761696	-74.001488	
729320	1	-74.006706	40.708244	-74.013550	
729321	4	-74.003342	40.743839	-73.945847	

	dropoff_latitude	store_and_fwd_flag	trip_duration
729317	40.789181	N	296
729318	40.731911	N	315
729319	40.741207	N	673
729320	40.713814	N	447
729321	40.712841	N	1224

```
[5]: data.shape
```

```
[5]: (729322, 11)
```

```
[6]: data.isnull().sum()
```

```
[6]: id                0
     vendor_id         0
     pickup_datetime   0
     dropoff_datetime  0
     passenger_count   0
     pickup_longitude   0
     pickup_latitude   0
     dropoff_longitude  0
     dropoff_latitude  0
     store_and_fwd_flag 0
```

```
trip_duration      0
dtype: int64
```

We have no missing values in this dataset.

```
[7]: data.dtypes
```

```
[7]: id                object
     vendor_id         int64
     pickup_datetime   object
     dropoff_datetime  object
     passenger_count   int64
     pickup_longitude   float64
     pickup_latitude   float64
     dropoff_longitude  float64
     dropoff_latitude  float64
     store_and_fwd_flag object
     trip_duration     int64
     dtype: object
```

```
[8]: # Checking for any duplicate records in the dataset
```

```
data.duplicated().sum()
```

```
[8]: 0
```

So, we have 0 duplicate records in the dataset which is good.

```
[9]: # Checking unique values in all columns
```

```
data.nunique()
```

```
[9]: id                729322
     vendor_id          2
     pickup_datetime   709359
     dropoff_datetime  709308
     passenger_count    9
     pickup_longitude   19729
     pickup_latitude   39776
     dropoff_longitude  27892
     dropoff_latitude  53579
     store_and_fwd_flag 2
     trip_duration     6296
     dtype: int64
```

Now, as per the dataset we have different types of variables which we will explore further each of them individually. But, we don't require the 'id' variable as it has 729322 unique values which we don't need for our model.

'trip\_duration' is the target variable and the values are continuous and hence this problem is a Regression Problem.

```
[10]: data['store_and_fwd_flag'].value_counts()
```

```
[10]: N    725282  
      Y     4040  
      Name: store_and_fwd_flag, dtype: int64
```

The store\_and\_fwd\_flag variable contains only two values i.e. Y and N. We can convert these values into numeric datatype by using Label Encoder.


```
[11]: encoder = LabelEncoder()  
  
      # Encoding column store_and_fwd_flag  
      data['store_and_fwd_flag_label'] = encoder.  
      ↪fit_transform(data['store_and_fwd_flag'])  
      data['store_and_fwd_flag_label'] = encoder.  
      ↪fit_transform(data['store_and_fwd_flag'])
```

```
[12]: # Converting pickup and dropoff from object to datetime  
  
      data['pickup_datetime'] = pd.to_datetime(data['pickup_datetime'])  
      data['dropoff_datetime'] = pd.to_datetime(data['dropoff_datetime'])
```

```
[13]: # Using datetime to create more new columns - day_number, pickup_hour, day_name  
  
      data['week_number'] = data.pickup_datetime.dt.weekday  
      data['pickup_hour'] = data.pickup_datetime.dt.hour  
      data['day'] = data.pickup_datetime.dt.day
```

```
[14]: # Plotting the pickup day wrt total no. of pickups  
  
      plt.figure(figsize = (22,6))  
      plt.subplot(131)  
      sns.countplot(data['day'])  
      plt.xlabel('Day')  
      plt.ylabel('Total No. of pickups')
```


```
[14]: Text(0, 0.5, 'Total No. of pickups')
```



```
[15]: # Plotting the hour wrt the total no. of pickups
```

```
plt.figure(figsize = (22,6))
plt.subplot(132)
sns.countplot(data['pickup_hour'])
plt.xlabel('Hour')
plt.ylabel('Total no. of pickups')
```


```
[15]: Text(0, 0.5, 'Total no. of pickups')
```



```
[16]: # plotting the week day wrt the total no. of pickups
```

```
plt.figure(figsize = (22,6))  
plt.subplot(133)  
sns.countplot(data['week_number'])  
plt.xlabel('Week day')  
plt.ylabel('Total no. of pickups')
```

```
[16]: Text(0, 0.5, 'Total no. of pickups')
```



#### 0.4.2 Insights:-

1. On Thursday(4), trip is on peak.
2. Trip is on peak between hour 17-22 i.e. in evenings of the day.
3. The no. of trip is high in the first three weeks.

```
[17]: data['drop_day'] = data.dropoff_datetime.dt.day
      data['drop_hour'] = data.dropoff_datetime.dt.hour
      data['drop_week'] = data.dropoff_datetime.dt.weekday
```

```
[18]: data.head()
```

```
[18]:
```

	id	vendor_id	pickup_datetime	dropoff_datetime	\
0	id1080784	2	2016-02-29 16:40:21	2016-02-29 16:47:01	
1	id0889885	1	2016-03-11 23:35:37	2016-03-11 23:53:57	
2	id0857912	2	2016-02-21 17:59:33	2016-02-21 18:26:48	



```

3 id3744273          2 2016-01-05 09:44:31 2016-01-05 10:03:32
4 id0232939          1 2016-02-17 06:42:23 2016-02-17 06:56:31

```

```

passenger_count pickup_longitude pickup_latitude dropoff_longitude \
0              1      -73.953918      40.778873      -73.963875
1              2      -73.988312      40.731743      -73.994751
2              2      -73.997314      40.721458      -73.948029
3              6      -73.961670      40.759720      -73.956779
4              1      -74.017120      40.708469      -73.988182

```

```

dropoff_latitude store_and_fwd_flag trip_duration \
0      40.771164              N           400
1      40.694931              N          1100
2      40.774918              N          1635
3      40.780628              N          1141
4      40.740631              N           848

```

```

store_and_fwd_flag_label week_number pickup_hour day drop_day \
0              0              0           16  29      29
1              0              4           23  11      11
2              0              6           17  21      21
3              0              1            9   5       5
4              0              2            6  17      17

```

```

drop_hour drop_week
0         16         0
1         23         4
2         18         6
3         10         1
4          6         2

```

```
[19]: # Getting distance calculated with the help of geopy
```

```

def calc_dist(data):
    pickup = (data['pickup_latitude'], data['pickup_longitude'])
    dropoff = (data['dropoff_latitude'], data['dropoff_longitude'])
    return geopy.distance.distance(pickup, dropoff).km

```

```
[20]: # Adding a new distance column
```

```
data['distance'] = data.apply(lambda x: calc_dist(x), axis=1)
```

```
[21]: data.head()
```

```

[21]:      id vendor_id pickup_datetime dropoff_datetime \
0 id1080784          2 2016-02-29 16:40:21 2016-02-29 16:47:01
1 id0889885          1 2016-03-11 23:35:37 2016-03-11 23:53:57

```

2	id0857912	2	2016-02-21 17:59:33	2016-02-21 18:26:48
3	id3744273	2	2016-01-05 09:44:31	2016-01-05 10:03:32
4	id0232939	1	2016-02-17 06:42:23	2016-02-17 06:56:31

	passenger_count	pickup_longitude	pickup_latitude	dropoff_longitude	\
0	1	-73.953918	40.778873	-73.963875	
1	2	-73.988312	40.731743	-73.994751	
2	2	-73.997314	40.721458	-73.948029	
3	6	-73.961670	40.759720	-73.956779	
4	1	-74.017120	40.708469	-73.988182	

	dropoff_latitude	store_and_fwd_flag	trip_duration	\
0	40.771164	N	400	
1	40.694931	N	1100	
2	40.774918	N	1635	
3	40.780628	N	1141	
4	40.740631	N	848	

	store_and_fwd_flag_label	week_number	pickup_hour	day	drop_day	\
0	0	0	16	29	29	
1	0	4	23	11	11	
2	0	6	17	21	21	
3	0	1	9	5	5	
4	0	2	6	17	17	

	drop_hour	drop_week	distance
0	16	0	1.199770
1	23	4	4.123945
2	18	6	7.250436
3	10	1	2.358287
4	6	2	4.328155

```
[22]: data['distance'].value_counts()
```

```
[22]: 0.000000    2901
      0.650138     2
      0.000424     2
      2.000050     1
      0.857289     1
      ...
      4.937193     1
      2.132841     1
      2.462624     1
      0.904627     1
      7.361812     1
      Name: distance, Length: 726420, dtype: int64
```

As per the dataset, there are 2901 records with distance travelled as 0.

```
[23]: data['passenger_count'].value_counts()
```

```
[23]: 1    517415
      2    105097
      5     38926
      3     29692
      6     24107
      4     14050
      0         33
      7          1
      9          1
      Name: passenger_count, dtype: int64
```

We have extremely low values for the passenger count above 6 and 0. We will remove these records.

```
[24]: # Considering only upto 6 passengers count

data = data[data['passenger_count'] <= 6]
data = data[data['passenger_count'] > 0]
```

```
[25]: # Again checking the passenger count

data['passenger_count'].value_counts()
```

```
[25]: 1    517415
      2    105097
      5     38926
      3     29692
      6     24107
      4     14050
      Name: passenger_count, dtype: int64
```

```
[26]: # Exploring the target variable - trip_duration
data['trip_duration'].describe()/3600
```

```
[26]: count    202.579722
      mean      0.264515
      std      1.073531
      min      0.000278
      25%      0.110278
      50%      0.184167
      75%      0.298611
      max      538.815556
      Name: trip_duration, dtype: float64
```

The max value in trip duration is 538 which is an outlier. We have to perform log transform to make more accurate predictions.

```
[27]: #trip duration in hours
data['trip_duration_hour'] = data['trip_duration'].apply(lambda x: x/3600)
```

```
[28]: #Removing outliers from the dataset such that the value does not exceeds a day,
      ↳ or 24 hours.
```

```
data = data[data['trip_duration_hour'] <= 24]
```

```
[29]: data['trip_duration_hour'].min()
```


```
[29]: 0.00027777777777777778
```

```
[30]: data['trip_duration_hour'].max()
```


```
[30]: 23.9975
```

```
[31]: sns.relplot(y = data['distance'], x = data['trip_duration_hour'])
```


```
[31]: <seaborn.axisgrid.FacetGrid at 0x2376d653be0>
```



```
[32]: sns.distplot(data['trip_duration'])  
plt.show()
```



```
[33]: # Performing log transform on trip duration in seconds  
  
data['trip_duration_log'] = np.log(data['trip_duration'].values + 1)  
sns.distplot(data['trip_duration_log'], kde = False)  
plt.show()
```



We are done with the Feature Exploration part, we can remove the unwanted columns from the dataset and even the outliers present.

```
[34]: df = data.drop(columns = ['id', 'pickup_datetime', 'dropoff_datetime',
    ↳ 'pickup_latitude', 'pickup_longitude',
    ↳ 'dropoff_latitude', 'dropoff_longitude',
    ↳ 'store_and_fwd_flag', 'trip_duration', 'trip_duration_hour'])
```

```
[35]: x = df.drop('trip_duration_log', axis=1)
y = df['trip_duration_log']
```

```
[36]: # Scaling the data

scaler = MinMaxScaler()
x_scale = scaler.fit_transform(x)

x = pd.DataFrame(x_scale, columns = x.columns)
```

As we have to predict discrete value and our target variable- trip\_duration is continuous, we can say that this is a Regression Problem.

**Evaluation Metric selected :- R-squared**

## 0.5 Building K Nearest Neighbor Model and finding the best value of k

```
[37]: train_x1, test_x1, train_y1, test_y1 = train_test_split(x, y, test_size = 0.25, random_state = 42)
```

```
[38]: # Elbow curve to determine the best value of k
def elbow(k):
    test = []

    for i in k:
        regr = KNeighborsRegressor(n_jobs = -1, n_neighbors = i)
        regr.fit(train_x1, train_y1)

        temp_pred = regr.predict(test_x1)
        temp_error = sqrt(mean_squared_error(temp_pred, test_y1)) # Using RMSE
        →for getting elbow curve
        test.append(temp_error)

    return test
```

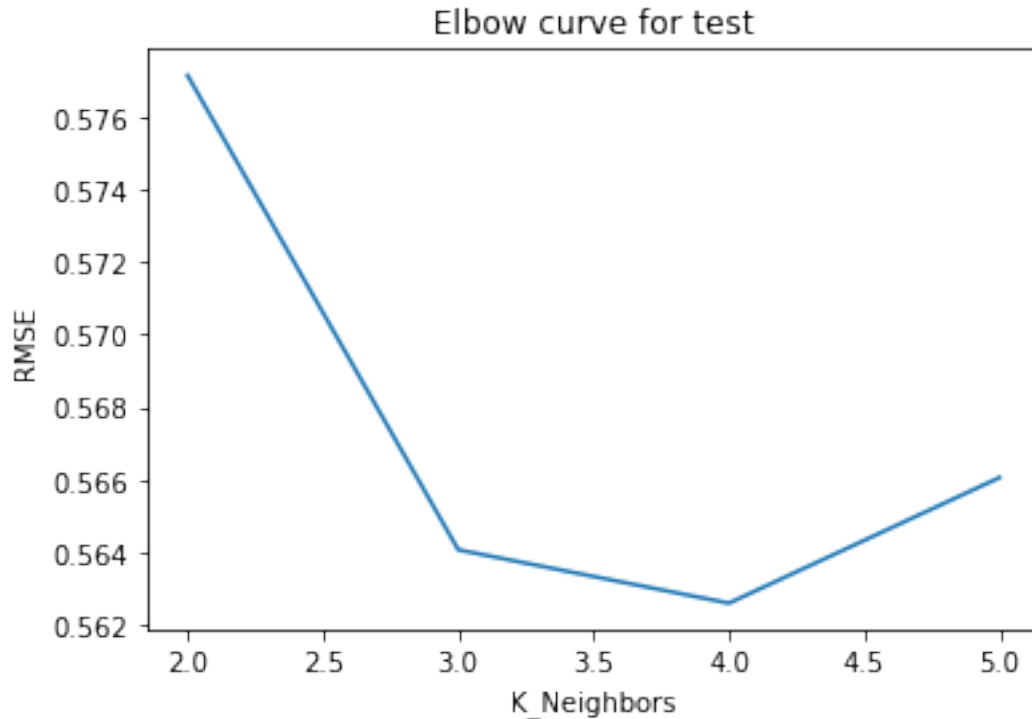
```
[39]: k = range(2, 6)
```

```
[40]: test1 = elbow(k)
```

```
[41]: # Plotting curve

plt.plot(k, test1)
plt.xlabel('K_Neighbors')
plt.ylabel('RMSE')
plt.title('Elbow curve for test')
```

```
[41]: Text(0.5, 1.0, 'Elbow curve for test')
```



As we can look at the elbow curve plot, the lowest error occurs at  $k=4$

```
[42]: knn = KNeighborsRegressor(n_jobs = -1, n_neighbors = 4)
      knn.fit(train_x1, train_y1)
```

```
[42]: KNeighborsRegressor(n_jobs=-1, n_neighbors=4)
```

```
[43]: # Prediction on test
      knn_pred = knn.predict(test_x1)

      # Checking R square score
      rscore = r2_score(test_y1, knn_pred)
      rscore
```

```
[43]: 0.5001540786344896
```

We got R-square score of KNN Model as - 0.50015

### 0.5.1 Building Linear Regression Model with Regularisation.

```
[44]: # Linear Model
      linear = LinearRegression()
      linear.fit(train_x1, train_y1)
```



```
[44]: LinearRegression()
```

```
[45]: pred = linear.predict(test_x1)

# Checking score
lscore = r2_score(test_y1, pred)
lscore
```

```
[45]: 0.36503803944085345
```

With Linear Regression, we got r-square score as - 0.36503

**Ridge Regression is a regularised linear regression model, So we use Ridge in this model.**

```
[46]: # Creating Ridge regression model
ridge_regr = Ridge(alpha = 1.0)
ridge_regr.fit(train_x1, train_y1)
```

```
[46]: Ridge()
```

```
[47]: # Getting Test score
ridge_pred = ridge_regr.predict(test_x1)

re_score = r2_score(test_y1, ridge_pred)
re_score
```

```
[47]: 0.34729058595313944
```

### 0.5.2 Building Random Forest model for the dataset.

```
[48]: random_for = RandomForestRegressor(random_state = 42)
random_for.fit(train_x1, train_y1)
```

```
[48]: RandomForestRegressor(random_state=42)
```

```
[49]: # Getting prediction and score

random_pred = random_for.predict(test_x1)

random_score = r2_score(test_y1, random_pred)
random_score
```

```
[49]: 0.6685600779144829
```

### 0.5.3 Building Gradient Boosting Model for the dataset.

```
[50]: xg_boost = XGBRegressor(random_state = 42, objective = 'reg:squarederror')
xg_boost.fit(train_x1, train_y1)
```

```
[50]: XGBRegressor(base_score=0.5, booster='gbtree', callbacks=None,
                  colsample_bylevel=1, colsample_bynode=1, colsample_bytree=1,
                  early_stopping_rounds=None, enable_categorical=False,
                  eval_metric=None, gamma=0, gpu_id=-1, grow_policy='depthwise',
                  importance_type=None, interaction_constraints='',
                  learning_rate=0.300000012, max_bin=256, max_cat_to_onehot=4,
                  max_delta_step=0, max_depth=6, max_leaves=0, min_child_weight=1,
                  missing=nan, monotone_constraints='()', n_estimators=100, n_jobs=0,
                  num_parallel_tree=1, predictor='auto', random_state=42,
                  reg_alpha=0, reg_lambda=1, ...)
```

```
[51]: # Getting prediction and score of the model
```

```
xg_pred = xg_boost.predict(test_x1)

xg_score = r2_score(test_y1, xg_pred)
xg_score
```

```
[51]: 0.7270941243460809
```

### 0.5.4 Combining all the different models to use averaging technique and generating final prediction.

```
[52]: # Averaging all the prediction from models- ridge, knn, random forest, xg boost,
      ↪ regressor and obtaining score.
```

```
combine_pred = (xg_pred + random_pred + ridge_pred + knn_pred)/4

combine_score = r2_score(test_y1, combine_pred)
combine_score
```

```
[52]: 0.6644762000679889
```

The R-square score we get by using averaging technique is:- 0.664476

```
[ ]:
```